

Efficient Large Language Model Deployment on Edge Devices

CSD415 Project Phase I Presentation

Izz Al Din Noufel Mukthar Joel K George Sarah Ann Pothen Shahma
Fathima

Department of Computer Engineering
Govt. Model Engineering College, Kochi

October 2025

Overview

1. Introduction
2. Problem Statement (LR)
3. Proposed Solution
4. Literature Survey (Detailed)
5. Objectives (Methodologies in LR)
6. Architecture
7. SRS
8. Conclusion
9. References
10. References

- Large Language Models (LLMs) such as GPT-4, PaLM, and LLaMA have revolutionized text understanding and generation.
- However, their massive scale (billions of parameters) prevents deployment on low-resource environments.
- Edge devices suffer from memory, compute, and power constraints.
- Bridging this gap requires compression pipelines that reduce model size while retaining linguistic and reasoning capabilities.

Problem Statement

- Cloud-based inference for LLMs incurs latency, privacy risks, and dependency on network availability.
- Running uncompressed models on-device leads to high inference delay and thermal throttling.
- Current compression methods often sacrifice accuracy or fail to generalize to multilingual and reasoning tasks.
- **Goal:** Enable practical and privacy-friendly on-device deployment of LLMs through an optimized hybrid compression framework.

Proposed Solution

- Build upon existing works in Knowledge Distillation (Acharya et al., Cantini et al.), Quantization (Lang et al.), and Pruning (Yi et al.).
- Develop a **hybrid compression pipeline** that:
 1. Distills semantic knowledge from large teacher models into smaller students.
 2. Applies structured pruning for attention heads and redundant layers.
 3. Employs mixed-precision quantization (8/4-bit) optimized for ARM and NPU hardware.
- Incorporate **Multi-Token Prediction (Samragh et al.)** to improve inference throughput.
- The approach aims to preserve up to 95% accuracy with an 80% reduction in model size.

Literature Survey Overview

- Extensive review conducted on 15+ recent works (2021–2025).
- Focused on pruning, quantization, and knowledge distillation (KD) techniques.
- Recent studies emphasize hybrid frameworks combining these methods for better trade-offs.
- Key trends:
 - **Symbolic Explainable KD** – improved interpretability.
 - **Hybrid Compression Pipelines** – joint pruning + quantization.
 - **Multi-token Inference** – improved throughput for LLMs.

Literature Survey Summary (2023–2025)

No.	Year	Name of Paper	Methodology	Advantage	Disadvantage
1	2024	Symbolic KD for LLMs (Acharya et al.)	Interpretable targets	Better interpretability	High complexity
2	2024	XAI-Driven KD for Low-Resource Devices (Cantini et al.)	Explainable KD	Transparent	Dependent on XAI quality
3	2023	Quantization Robust Pruning (Kim)	Pruning + KD	High compression	Complex tuning
4	2023	Mutual Learning Framework (JCST Team)	Multi-network KD	Stable training	Resource overhead
5	2023	Efficient Hybrid Compression (Malihi et al.)	KD + Pruning + Fine-Tuning	Accuracy control	Depends on fine-tuning
6	2023	Quantization Techniques for LLMs (Lang et al.)	Comparative study	Hardware insights	Accuracy drop risk
7	2023	Advances in Pruning and Quantization (Bibi et al.)	Survey	Broad coverage	Limited LLM data

Table: Key Works in LLM Compression (Part I)

Literature Survey Summary (2023–2025 Continued)

No.	Year	Name of Paper	Methodology	Advantage	Disadvantage
8	2023	Linear Quantization Optimization (Algorithms Team)	Low-bit quantization	Efficient edge use	Hardware dependency
9	2023	Federated Layer-wise Pruning (Entropy Team)	Layer pruning in FL	Adaptive edge design	Comm. trade-offs
10	2025	EdgeMoE (Yi et al.)	Sparse mixture-of-experts	High sparsity	Loss of expressiveness
11	2024	Adaptive Scaling Pruning (Ko, Kang)	Filter pruning for vision	Efficient edge vision	Domain-specific
12	2024	Hybrid Weighted Pruning (Geng et al.)	Weighted hybrid pruning	High FLOP reduction	Fine-tuning needed
13	2023	Statistical Pruning for IoT (Mao et al.)	Statistical pruning	Low power usage	IoT-specific
14	2021	Sparsity in Deep Learning (Hoeffler et al.)	Growth-pruning survey	Contextual pruning insight	No LLM focus
15	2025	Multi-Token Prediction (Samragh et al.)	Multi-token decoding	Faster inference	Experimental

Table: Key Works in LLM Compression (Part II)


Objectives

- Design a **modular compression framework** integrating Knowledge Distillation (KD), pruning, and quantization.
- Evaluate **compression-accuracy trade-offs** for different LLM architectures.
- Develop an **adaptive runtime** for efficient inference across diverse edge hardware.
- Establish **reproducible pipelines** with defined benchmarks for latency, throughput, and accuracy.

Expected Outcomes

- **Efficiency:** 80% model compression, 2–3× faster inference.
- **Scalability:** Portable across CPU, GPU, and NPU.
- **Interpretability:** Distilled models retain explainable attention patterns.
- **Edge-readiness:** Suitable for on-device AI assistants and IoT endpoints.

System Architecture



Software Requirements Specification

- **Functional Requirements:**

- Support KD, pruning, and quantization modules.
- APIs for compression, deployment, and profiling.

- **Non-Functional Requirements:**

- Latency $\leq$ 150 ms (first token)
 - Throughput $\geq$ 20 tokens/s
 - Memory footprint $\leq$ 1.5 GB
- Design documents, SRS, and architecture blueprints have been prepared to formalize implementation.

System Design Overview (Part I)

- Combines symbolic KD (Acharya et al.) with quantization-aware training (Lang et al.).
- Supports structured pruning guided by model saliency maps.
- Utilizes a flexible recipe-based configuration for hardware adaptation.
- Edge runtime supports CPU, GPU, and NPU backends with dynamic precision scaling.




Figure: Overall System Architecture

System Design Overview (Part II)




Figure: Edge Runtime Flow Diagram

Conclusion

- The study focused on enabling efficient LLM deployment on resource-constrained environments.
- Existing literature demonstrated strong potential in hybrid compression methods.
- Building on these insights, our project designs a practical pipeline combining KD, pruning, and quantization.
- We have completed SRS, design documentation, and architectural planning for the proposed model.
- Future phases will focus on implementation, benchmarking, and real device testing.

References I

-  K. Acharya et al., *"A Survey on Symbolic Knowledge Distillation of Large Language Models,"* IEEE Transactions on AI, 2024. <https://arxiv.org/pdf/2408.10210>
-  R. Cantini et al., *"XAI-Driven Knowledge Distillation of LLMs for Low-Resource Devices,"* Journal of Big Data, 2024. <https://journalofbigdata.springeropen.com/counter/pdf/10.1186/s40537-024-00928-3.pdf>
-  J. Kim, *"Quantization Robust Pruning with Knowledge Distillation,"* IEEE Access, 2023. <https://ieeexplore.ieee.org/stamp/stamp.jsp?tp=&arnumber=10072406>
-  L. Malihi et al., *"Efficient and Controllable Model Compression through Knowledge Distillation, Pruning, and Fine-Tuning,"* Neurocomputing, 2023.
-  J. Lang et al., *"A Comprehensive Study on Quantization Techniques for Large Language Models,"* ACM Computing Surveys, 2023.
-  U. Bibi et al., *"Advances in Pruning and Quantization for NLP,"* Journal of AI Research, 2023.

References II

-  R. Yi et al., “*EdgeMoE: Empowering Sparse Large Language Models for Edge Deployment*,” IEEE Transactions on Mobile Computing, 2025.
<https://ieeexplore.ieee.org/stamp/stamp.jsp?arnumber=10664567>
-  X. Geng et al., “*Complex Hybrid Weighted Pruning Method for Accelerating Neural Networks*,” IEEE Transactions on Neural Networks and Learning Systems, 2024.
-  M. Samragh et al., “*Your LLM Knows the Future: Uncovering Its Multi-Token Prediction Potential*,” arXiv:2507.11851, 2025. <https://arxiv.org/pdf/2507.11851>
-  T. Hoefler et al., “*Sparsity in Deep Learning: Pruning and Growth for Efficient Inference*,” ACM Computing Surveys, 2021.

Thank You!